

## 01. Introduction

- **Agent** - Anything that can perceive its environment through sensors and acting upon that env. through actuators
- **Agent Function** - Maps from percept histories to actions
- **Rational Agent** - Chooses an action that is expected to maximize its performance measure, given by percept sequence and built-in knowledge
- **Autonomous Agent** - If behavior is determined by its own experience

### Performance Measure of Function

- Motivation: For an agent to do the right thing, need a measure of goodness
- Performance vs. Cost
  1. Best for whom?
  2. What are we optimizing?
  3. What information is available?
  4. What are the side effects and costs?

### Defining the Problem: PEAS

1. Performance measure
2. Environment
3. Actuators
4. Sensors

### Characterizing the Environment

1. **Fully observable** - (vs. Partially) Agent's sensors can access complete state of env. all the time
2. **Deterministic** - (vs. Stochastic) Next state of env. is determined by **current state** and **action executed by agent**
  - **Strategic** - If env. is deterministic except for actions of other agents
3. **Episodic** - (vs. Sequential) Agent's experience is divided into atomic **episodes**, where each episode includes perceiving and an action, and **action depends on episode** itself
4. **Static** - (vs. Dynamic) Env. is unchanged while agent is deciding
  - **Semi** - Time does not affect env., but affects performance score
5. **Discrete** - (vs. Continuous) Discrete num. of percepts and actions
6. **Single Agent** - (vs. Multi-agent) Agent operating by itself in an env.

### Implementing Agents (in ascending complexity)

1. **Simple Reflex Agents** - Fixed conditional rules
2. **Model-based Reflex Agents** - Stores percept history to make decisions about internal model of world with conditional rules. Eg. Roomba
3. **Goal-based Agents** - Keep in mind a goal and action aims to achieve it
4. **Utility-based Agents** - Find best way to achieve goal
5. **Learning Agents** - Learn from previous experiences

### Exploitation vs. Exploration

- **Exploitation** - Maximize expected utility using current knowledge of world
- **Exploration** - Learn more about the world to improve future gains. May not always maximize performance measure.

## 02. Uninformed Search

- Deterministic, fully observable
- **Tree Search** - Can revisit nodes
- **Graph Search** - Tracks visited (Tree Search + Memoization)
- **Uninformed Search** - Uses only information available in problem definition

### Formulating the Problem

1. How to represent state in problem?
2. Initial state
3. Actions: Successor function
4. Goal test
5. Path cost

- **Abstraction Function** - Maps abstracted representation to real world state
- **Representation Invariant** -  $I(c) = \text{True} \rightarrow \exists a \text{ s.t. } AF(c) = a$

### Breadth-first Search

- Idea: Expand shallowest unexpanded node using **queue**
- Given:  $b$ : Branching factor and  $d$ : Depth of optimal solution
- Complete: Yes (if tree is finite)
- Time:  $O(b^{d+1})$
- Space:  $O(b^d)$
- Optimal: Yes (if cost = 1)
- BFS is Uniform-cost Search with same cost

### Uniform-cost Search

- Idea: Expand least-cost unexpanded node using **priority queue** (Dijkstra's)
- Given:  $C^*$ : Cost of optimal solution
- Complete: Yes (if step cost  $\geq \epsilon$  where  $\epsilon \geq 0$ )
- Time:  $O(b^{(C^*/\epsilon)})$  ( $C^*/\epsilon$  is approx. number of layers)
- Space:  $O(b^{(C^*/\epsilon)})$
- Optimal: Yes

### Depth-first Search

- Idea: Expand deepest unexpanded node using **stack**
- Given:  $m$ : Maximum depth of tree
- Complete: No (fails with infinite depth or loops)
- Time:  $O(b^m)$
- Space:  $O(bm)$  (better than BFS)
- Optimal: No

### Depth-limited Search

- DFS with depth limit  $l$  where nodes at depth  $l$  have no children
- Time:  $b^0 + b^1 + \dots + b^{(d-1)} + b^d = O(b^d)$

### Iterative Deepening Search

- Idea: Try different depths for depth-limited search
- Complete: Yes
- Time:  $(d+1)b^0 + db^1 + \dots + b^d = O(b^d)$  (More overhead than DLS)
- Space:  $O(bd)$

### Summary

	BFS	Uniform Cost	DFS	DLS	IDS
<b>Complete</b>	Yes	Yes	No	No	Yes
<b>Time</b>	$O(b^d)$	$O(b^{C^*/\epsilon})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
<b>Space</b>	$O(b^d)$	$O(b^{C^*/\epsilon})$	$O(bm)$	$O(bl)$	$O(bd)$
<b>Optimal</b>	Yes	Yes	No	No	No

### Bidirectional Search

- Idea: Search both forwards from initial state and backwards from goal state. Stop when searches meet.
- Time:  $O(2b^{d/2})$
- Operators must be reversible
- Can have many goal states
- How to check if node intersects with other half?

## 03. Informed Search

## 04. Adversarial Search

- Assumption: Opponents reacts rationally
- Needs:
  - Initial state
  - Successor function
  - Terminal test
  - Utility function: Measures how good the move is for a player

### Minimax

- Idea: Choose move that yields highest minimax value

## 05. Introduction to Machine Learning

- A machine learns if it improves performance P on task T based on experience E. Where T must be fixed, P must be measurable, E must exist

### Types of Feedback

- **Supervised** - Correct answer given for each example
  - **Regression** - Predict results within continuous output
  - **Classification** - Predict results in discrete output
- **Unsupervised** - No answers given
- **Weakly supervised** - Answer given, but not precise
- **Reinforcement** - Occasional rewards given

### Decision Trees


- DT can express any function of input attributes, if data is consistent
- Goal: Make DT **compact**. How?

Information Theory

- Idea: Choose attribute that splits examples into subsets that are ideally 'all positive' or 'all negative'
- Entropy** - Measure of randomness in set of data

$$I(P(v_1), ..., P(v_n)) = - \sum_{i=1}^n P(v_i) \log_2 P(v_i)$$

- For data with *p* positive examples and *n* negative examples:



- Information Gain** - (IG) Reduction in entropy from attribute test
- Goal: Choose attribute with largest information gain
- Intuition: IG = Entropy of this node - Entropy of children nodes
- Given chosen attribute *A* with *v* distinct values:

$$\text{remainder}(A) = \sum_{i=1}^v \frac{p_i + n_i}{p + n} I(\frac{p_i}{p_i + n_i}, \frac{n_i}{p_i + n_i})$$
$$IG(A) = I(\frac{p}{p+n}, \frac{n}{p+n}) - \text{remainder}(A)$$

- Decision Tree Learning** - Recursively choose attributes with highest IG
- IG is not the only way. Can use whatever objective function that achieves the criteria we want.

Performance Measurement

- Correctness** - Correct if *ŷ* = *y*
- Accuracy** -  $\frac{1}{m} \sum_{j=1}^m (\hat{y}_j = y_j)$
- Confusion Matrix:

		Actual Label	
		+ve	-ve
Predicted Label	+ve	TP True Positive	FP False Positive
	-ve	FN False Negative	TN True Negative

- Accuracy =  $\frac{TP+TN}{TP+FN+FP+TN}$
- Precision** -  $\frac{TP}{TP+FP}$  **Recall** -  $\frac{TP}{TP+FN}$
- Type I Error: FP Type II Error: FN
- FP Rate =  $\frac{FP}{FP+TN}$  TP Rate =  $\frac{TP}{TP+FN}$

Pruning

- Motivation: DT overfits to training set, but performs poorly on test set
- Occam's Razor: Simple hypothesis preferred
- Pruning** - Ignores outliers, which reduces overfitting
  - E.g. Min-sample, Max-depth

06. Linear Regression

Notation

- m* = Number of training examples
- n* = Number of features
- x<sub>j</sub><sup>(i)</sup>* = Input feature *j* of *i*th training example
- y* = Output variables

Hypothesis

$$h_w(x) : w_0 + w_1x$$

Cost Function (Square Error Function)

$$J(w_0, w_1) = \frac{1}{2m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)})^2$$

- Goal: Minimize cost function. Thus, hypothesis is close to training samples
- Why squared error? Convenience, since we need to differentiate later

Gradient Descent

- Start at some (*w*<sub>0</sub>, *w*<sub>1</sub>). Pick nearby point that reduces *J*(*w*<sub>0</sub>, *w*<sub>1</sub>).
- Algorithm: Repeat until convergence:

$$w_j := w_j - \alpha \frac{dJ(w_0, w_1, \dots)}{dw_j}$$

- All updates done at end
- How to do  $\frac{dJ(w_0, w_1)}{dw_j}$ ? Partial derivative: Hold everything else constant
  - $\frac{dJ(w_0, w_1)}{dw_j} = \frac{d}{dw_j} (\frac{1}{2m} \sum_{i=1}^m (w_0 + w_1x^{(i)} - y^{(i)})^2)$
  - $\frac{dJ(w_0, w_1)}{dw_0} = \frac{1}{m} \sum_{i=1}^m (w_0 + w_1x^{(i)} - y^{(i)})$  (Note: Chain rule)
  - $\frac{dJ(w_0, w_1)}{dw_1} = \frac{1}{m} \sum_{i=1}^m (w_0 + w_1x^{(i)} - y^{(i)})x^{(i)}$
- Time complexity: *O*(*kmn*) where *k* is number of iterations

Learning Rate

- If *α* too small, then descent is too slow. If *α* too big, then might overshoot.
- Given constant *α*, descent will grow smaller as we approach minimum

Variants of Gradient Descent

- Batch gradient descent: Consider all training examples when updating
- Stochastic gradient descent: Consider 1 random data point at a time (Cheaper and more randomness)
- Mini-batch gradient descent

Using Matrices

- Given:  $w = \begin{pmatrix} w_0 \\ \vdots \\ w_n \end{pmatrix}$  and  $x = \begin{pmatrix} x_0 \\ \vdots \\ x_n \end{pmatrix} = \begin{pmatrix} 1 \\ \vdots \\ x_n \end{pmatrix}$
- $h_w(x) : w^T x$

Feature Scaling

- Motivation: Gradient descent does not work well if features have different scales
- Mean Normalization** -  $x_i \leftarrow \frac{x_i - \mu_i}{\sigma_i}$

Normal Equation

$$w = (X^T X)^{-1} X^T Y$$

- No need to choose *α* and feature scaling
- X<sup>T</sup>X* needs to be invertible
- Time complexity: *O*(*n*<sup>3</sup>). Slow if *n* is big

07. Logistic Regression