

## 01. Introduction

- $f(x, y) = i(x, y)r(x, y)$  where  $f$  is intensity,  $i$  is illumination, and  $r$  is reflectance
- Exposure time** - Time for incident light to reach sensor
- Storage** -  $b = MNK$  where img. has size  $M \times N$  and  $K$ -bit depth
- Color to greyscale:  $I = W_R R + W_G G + W_B B$  where  $\sum_{i \in \{R, G, B\}} W_i = 1$
- Common weights:  $(.299, .587, .114)$
- Norm. RGB** -  $(r, g, b) = (\frac{R}{R+G+B}, \frac{G}{R+G+B}, \frac{B}{R+G+B})$
- $(R, G, B) \leftrightarrow (r, g, b) \leftrightarrow (r, g, I)$  where  $I = \frac{R+G+B}{3}$
- HSV Color Space**
  - Hue: Pure color (0 to 360)
  - Saturation: Mix pure color (1) with white light (0)
  - Value: Mix from black (0) to white (255)

## 02. Filtering

- Point processing** -  $x_{ij} = f(p_{ij})$ 
  - Brightness** -  $x_{ij} = p_{ij} + b$  (Clipping behavior)
  - Intensity scaling** -  $x_{ij} = ap_{ij}$  (Increased/decreased contrast)
- Normalization (Whitening)** -  $x_{ij} = \frac{p_{ij} - \mu}{\sigma} \sigma^2 = \frac{\sum_{i=0}^I \sum_{j=0}^J (p_{ij} - \mu)^2}{IJ}$
- Gamma** -  $x_{ij} = 255(\frac{p_{ij}}{255})^\gamma$  where  $\gamma > 0$
- Intensity Histogram** - No data on location
  - Stretching** -  $x = (p - f_{\min})(\frac{255}{f_{\max} - f_{\min}})$
  - Equalization** - Turn cumulative distribution linear
    - Get histogram:  $h_k = \sum_{i=0}^I \sum_{j=0}^J \delta$  where  $\delta = 1$  if  $p_{ij} - k = 0$  and 0 otherwise
    - Estimate CDF:  $c_k = \frac{\sum_{l=1}^k h_l}{\sum_{l=1}^J h_l}$
    - Map  $p_{ij} = k$  to new bin  $x_{ij}$ :  $x_{ij} = 255 \text{CDF}(k)$
    - Pros: More contrast than normalization, less prone to outliers than stretching
    - Cons: More expensive
- Good for foreground and background separation: Bi-modal histogram can be thresholded
- Otsu's Method** - Automated thresholding
  - $T^* = \min_{T \in [0, 255]} (w_1(T)\sigma_1^2(T) + w_2(T)\sigma_2^2(T))$  where  $T^*$  is the optimal threshold that min. sum of weighted variances of object and background
  - $\sigma_1^2(T)$  and  $\sigma_2^2(T)$ : Variance of pixels less than or equal to and greater than threshold respectively
  - $w_1(T)$  and  $w_2(T)$ : Number of pixels less than or equal to and greater than threshold
- Correlation Filtering** - Window (Moving average)
  - Motivation: Reduce noise

- Impulse noise: Random white occurrences of pixels
- Salt and pepper noise: Rand. white and black pixels
- Gaussian noise: Variations in intensity from Gaussian distribution  $p_{ij} = \hat{p}_{ij} + \eta$  where  $\eta \sim N(\mu_n, \sigma_n)$

- $x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$  where  $f$  is a **kernel** of weights
- Notation:  $X = P \otimes F$

- Box Filter** - Blurs image

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Given kernel of width  $2k+1$ , add padding:  $\lfloor \frac{2k+1}{2} \rfloor$
- Filling methods: Zero-padding, wrap around, copy edge, reflect
- Gaussian Filter** - Nearest neighboring pixels have more weight

$$f_{uv} = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

- Size of kernel: At some threshold, kernel will have lots of 0s, which is not useful
- Variance: Determines smoothing effect

- Template Matching** - Set template as kernel and match occurs at local max.

- Normalized Cross-Correlation** - Normalizes using filter and input window
- Motivation: Non-normalized is dominated by original image pixels

- $x_{ij} = \frac{1}{|F||w_{ij}|} \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$  where  $|\cdot|$  is magnitude of kernel (i.e. Square root of sum of squares)

- Convolution Filtering** - Flip kernel in both directions, then do cross-correlation

$$x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i-u, j-v}$$

- Notation:  $X = F * P$

- Convolution has nicer properties than correlation

- Properties: Commutative, Associative, Distributive over addition, Scalar factor, Identity

- Non-Linear Filters** - Filters that perform non-linear operations (e.g. Median, Min., Max.)

- Correlation and convolution are both linear operations

- Median Filtering**

- Removes spikes: Good for removing impulse and salt and pepper noise
- No new pixel values introduced
- Edge preserving

## 03. Gradients

- Gradient** - Vector at pixel pointing at most rapid change in intensity

$$\nabla f = \langle \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \rangle$$

- In images, der. is finite:  $f'(x) = \frac{f(x+1) - f(x-1)}{2}$

$$\theta = \tan^{-1}(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x})$$

$$\|\nabla f\| = \sqrt{(\frac{\partial f}{\partial x})^2 + (\frac{\partial f}{\partial y})^2}$$

- Sobel Filter**

- Horizontal sobel filter ( $S_x$ ): Detects vertical lines
- Vertical sobel filter ( $S_y$ ): Detects horizontal lines

$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \quad S_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- Differentiation is sensitive to noise: Need to blur first

- Trade-off: Blur removes noise, but blurs the edge
- Derivative of Gaussian** - Blur + Derivative merged into 1 filter;  $\frac{\partial}{\partial x}(h * f) = (\frac{\partial}{\partial x}h) * f$

- Laplace Filter** - 2nd derivative filter

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - 2f(x) + f(x-h)}{h^2}$$


- Laplacian of Gaussian (LoG) Filter** - Combine Laplace filter with Gaussian
  - Produces zero-crossings at edge

- Canny Edge Detector**

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Thin wide ridges down to single pixel width via non-maximum suppression
- Perform thresholding (hysteresis) and linking

- Non-maximum Suppression** - For each pixel, check if it's local maximum along gradient direction. If yes, then keep. Else, 0.

- Interpolation:




- Alternative: Approx. using closest pixel value

- Hysteresis Thresholding** - High threshold to start edge curves and low threshold to continue growing

- Motivation: Single high threshold can miss out some parts of the edge

04. Lines

- Goal: Which points belong to which line?
- $y = mx + b \quad \frac{x}{a} + \frac{y}{b} = 1$
- $x \cos \theta + y \sin \theta = \rho$
- **Line fitting** -  $E = \frac{1}{N} \sum_i (y_i - (mx_i + b))^2$ 
  - Goal: Find parameters  $m$  and  $b$  that min.  $E$
  - Using gradient descent:  $b = \bar{y} - m\bar{x} \quad m = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$  where  $\bar{y} = \frac{\sum_i y_i}{N}$  and  $\bar{x} = \frac{\sum_i x_i}{N}$
  - Cons: Choice of error function, outliers
- **Hough Transform** - Paramterize shape and vote



1. Initialize accumulator array  $A(\theta, \rho) = 0$
  2. For each image edge point  $(x_i, y_i)$ :
    1. For each  $\theta$ :
      1. Solve for  $\rho = x_i \cos \theta + y_i \sin \theta$
      2.  $A(\theta, \rho) = A(\theta, \rho) + 1$
  3. Find maximum in  $A(\theta, \rho)$
  4. Detected lines are given by  $\rho^* = x \cos \theta^* + y \sin \theta^*$
- **Circle** -  $(x - a)^2 + (y - b)^2 = r^2$ 
    - If  $r$  is unknown, need to solve for 3 parameters (i.e.  $A(a, b) \rightarrow A(a, b, r)$ ) and a point becomes a cone in 3D parameter space
    - Optimization: Use gradient
      - Use edge orientation  $\phi_i$  to vote for 2 points, rather than whole circle
  - **Generalized Hough Transform** - Arbitrary shape
    1. Given shape with boundary points  $p_i$  and reference point  $a$
    2. For each  $p_i$ , get displacement vector from  $a$
    3. Store in table with key  $\phi_i$  and value (Displacement vectors)
    4. For each edge pt., use  $\phi_i$  to get vectors to vote for  $a$
    - Application: Index with local pattern, instead of gradients
    - Tips: Soft voting, convert to edge image
    - Pros: Food with noise and occlusion
    - Cons: Complexity with num. of parameters, grid size

05. Segmentation

- Goal: Separate image into coherent regions
- Idea: **Clustering** - Group similar data points together
- Challenges: What makes 2 points same/different? Choice of features (e.g. Color, Intensity, Position), Which clustering algorithm?
- **k-Means Clustering** - Iteratively re-assign points to nearest cluster center
  1. Randomly initialize the cluster centers  $c_1, \dots, c_K$

2. For each point  $p_i$ , find the closest  $c_j$  to put  $p_i$  in
3. Set  $c_j$  to be mean of points in cluster  $j$
4. Repeat, if  $c_j$  have changed up to some threshold

- Pros: Simple, Converges to local min.
- Cons: Setting  $K$ , Sensitive to initial centers (Since k-means converges to local min.), Sensitive to outliers (Can add more clusters), Assumes spherical clusters

• **Simple Linear Iterative Clustering (SLIC) Superpixels**

- **Superpixel** - Group of pixels that share common traits
  - Application: Inputs to other CV algo. since more compact representation with perceptual meaning
- Num. of pixels:  $n_{tp}$ ; Target num. of superpixels:  $n_{sp}$
- Initial width of each superpixel:  $s = \sqrt{n_{tp}/n_{sp}}$
- Features:  $z = [r, g, b, x, y]$
- Color distance:  $d_c = ||\langle r_j, g_j, b_j \rangle - \langle r_i, g_i, b_i \rangle||$
- Spatial distance:  $d_s = ||\langle x_j, y_j \rangle - \langle x_i, y_i \rangle||$
- Scaling factors:  $d_{cm}$  and  $d_{sm} = s$  set as max. expected values of  $d_c$  and  $d_s$  respectively
- $D = \sqrt{(\frac{d_c}{d_{cm}})^2 + (\frac{d_s}{d_{sm}})^2} = \sqrt{d_c^2 + (\frac{d_s}{s})^2} c^2$

1. Split img. into grid of size  $s \times s$ . Set cluster centers as lowest gradient position in  $3 \times 3$  neighborhood from superpixel center to speed up convergence since initialize on value common to surrounding.
  2. For each cluster center, check distance to all pixels within  $2s \times 2s$  neighborhood. Assign pixels to closest checked center.
  3. Update cluster centers using mean and repeat if not converged (Same as k-Means)
  4. Optional: Replace superpixel region by average value to create stained glass effect
- Modification of k-Means: Not random initialization, Compute pixel's distance only to closest set of cluster centers

• **Mean-Shift Clustering** - Find local density maxima in feature space

- **Attraction basin** - Region in feature space for which all trajectories of centroids lead to same mode
- **Cluster** - All data points in attraction basin of a mode

1. For each data point:
  1. Define window around and get centroid
  2. Shift window to centroid
  3. Repeat until window centroid stops moving

- Segmentation with Mean Shift: Do mean shift and merge pixels in same attraction basin

- Choosing window size: Trial and error, Sample points and use avg. dist. to knn. (Num. of neighbors needs to be large enough to ensure increase in density)

- Larger window size  $\rightarrow$  Fewer clusters

- Pros: No assumptions on cluster shape, 1 parameter, Finds variable num. of modes (vs. specified  $k$  in k-Means), Robust to outliers
- Cons: Choosing  $h$ , Slow, Scales poorly with feature space dimension
- Optimizations:
  - After each run of mean shift, assign all points within radius  $r$  of end point to same cluster
  - Assign points in radius  $c < r$  of search path to mode