

01. Introduction

- $f(x, y) = i(x, y)r(x, y)$ where f is intensity, i is illumination, and r is reflectance
- **Exposure time** - Time for incident light to reach sensor
- **Storage** - $b = MNK$ where img. has size $M \times N$ and K -bit depth
- Color to greyscale: $I = W_R R + W_G G + W_B B$ where $\sum_{i \in [R, G, B]} W_i = 1$
 - Common weights: (.299, .587, .114)
- **Norm. RGB** - $(r, g, b) = (\frac{R}{R+G+B}, \frac{G}{R+G+B}, \frac{B}{R+G+B})$
- $(R, G, B) \leftrightarrow (r, g, b) \leftrightarrow (r, g, I)$ where $I = \frac{R+G+B}{3}$
- **HSV Color Space**
 - Hue: Pure color (0 to 360)
 - Saturation: Mix pure color (1) with white light (0)
 - Value: Mix from black (0) to white (255)

02. Filtering

- **Point processing** - $x_{ij} = f(p_{ij})$
 - **Brightness** - $x_{ij} = p_{ij} + b$ (Clipping behavior)
 - **Intensity scaling** - $x_{ij} = ap_{ij}$ (Increased/decreased contrast)
 - **Normalization (Whitening)** - $x_{ij} = \frac{p_{ij} - \mu}{\sigma} \sigma^2 = \frac{\sum_{i=0}^I \sum_{j=0}^J (p_{ij} - \mu)^2}{IJ}$
 - **Gamma** - $x_{ij} = 255(\frac{p_{ij}}{255})^\gamma$ where $\gamma > 0$
 - **Intensity Histogram** - No data on location
 - **Stretching** - $x = (p - f_{\min})(\frac{255}{f_{\max} - f_{\min}})$
 - **Equalization** - Turn cumulative distribution linear
 1. Get histogram: $h_k = \sum_{i=0}^I \sum_{j=0}^J \delta$ where $\delta = 1$ if $p_{ij} - k = 0$ and 0 otherwise
 2. Estimate CDF: $c_k = \frac{\sum_{l=1}^k h_l}{IJ}$
 3. Map $p_{ij} = k$ to new bin x_{ij} : $x_{ij} = 255 \text{CDF}(k)$
 - Pros: More contrast than normalization, less prone to outliers than stretching
 - Cons: More expensive

03. Gradients

04. Lines

05. Segmentation

- Goal: Separate image into coherent regions
- Idea: **Clustering** - Group similar data points together
- Challenges: What makes 2 points same/different? Choice of features (e.g. Color, Intensity, Position), Which clustering algorithm?
- **k-Means Clustering** - Iteratively re-assign points to nearest cluster center
 1. Randomly initialize the cluster centers c_1, \dots, c_K
 2. For each point p_i , find the closest c_j to put p_i in
 3. Set c_j to be mean of points in cluster j
 4. Repeat, if c_j have changed up to some threshold
 - Pros: Simple, Converges to local min.

- Cons: Setting K , Sensitive to initial centers (Since k-means converges to local min.), Sensitive to outliers (Can add more clusters), Assumes spherical clusters
- **Simple Linear Iterative Clustering (SLIC) Superpixels**
 - **Superpixel** - Group of pixels that share common traits
 - Application: Inputs to other CV algo. since more compact representation with perceptual meaning
 - Num. of pixels: n_{tp} ; Target num. of superpixels: n_{sp}
 - Initial width of each superpixel: $s = \sqrt{n_{tp}/n_{sp}}$
 - Features: $z = [r, g, b, x, y]$
 - Color distance: $d_c = ||\langle r_j, g_j, b_j \rangle - \langle r_i, g_i, b_i \rangle||$
 - Spatial distance: $d_s = ||\langle x_j, y_j \rangle - \langle x_i, y_i \rangle||$
 - Scaling factors: d_{cm} and $d_{sm} = s$ set as max. expected values of d_c and d_s respectively
 - $D = \sqrt{(\frac{d_c}{d_{cm}})^2 + (\frac{d_s}{d_{sm}})^2} = \sqrt{d_c^2 + (\frac{d_s}{s})^2} c^2$
- 1. Split img. into grid of size $s \times s$. Set cluster centers as lowest gradient position in 3×3 neighborhood from superpixel center to speed up convergence since initialize on value common to surrounding.
- 2. For each cluster center, check distance to all pixels within $2s \times 2s$ neighborhood. Assign pixels to closest checked center.
- 3. Update cluster centers using mean and repeat if not converged (Same as k-Means)
- 4. Optional: Replace superpixel region by average value to create stained glass effect
 - Modification of k-Means: Not random initialization, Compute pixel's distance only to closest set of cluster centers
- **Mean-Shift Clustering** - Find local density maxima in feature space
 - **Attraction basin** - Region in feature space for which all trajectories of centroids lead to same mode
 - **Cluster** - All data points in attraction basin of a mode
- 1. For each data point:
 - (1) Define window around and get centroid
 - (2) Shift window to centroid
 - (3) Repeat until window centroid stops moving
- Segmentation with Mean Shift: Do mean shift and merge pixels in same attraction basin
- Choosing window size: Trial and error, Sample points and use avg. dist. to knn. (Num. of neighbors needs to be large enough to ensure increase in density)
 - Larger window size \rightarrow Fewer clusters
- Pros: No assumptions on cluster shape, 1 parameter, Finds variable num. of modes (vs. specified k in k-Means), Robust to outliers
- Cons: Choosing h , Slow, Scales poorly with feature space dimension
- Optimizations:
 - After each run of mean shift, assign all points within radius r of end point to same cluster
 - Assign points in radius $c < r$ of search path to mode