

01. Introduction

- $f(x, y) = i(x, y)r(x, y)$ where f is intensity, i is illumination, and r is reflectance
- Exposure time** - Time for incident light to reach sensor
- Storage** - $b = MNK$ where img. has size $M \times N$ and K -bit depth
- Color to greyscale: $I = W_R R + W_G G + W_B B$ where $\sum_{i \in [R, G, B]} W_i = 1$
 - Common weights: (.299, .587, .114)
- Norm. RGB** - $(r, g, b) = (\frac{R}{R+G+B}, \frac{G}{R+G+B}, \frac{B}{R+G+B})$
- $(R, G, B) \leftrightarrow (r, g, b) \leftrightarrow (r, g, I)$ where $I = \frac{R+G+B}{3}$
- HSV Color Space**
 - Hue: Pure color (0 to 360)
 - Saturation: Mix pure color (1) with white light (0)
 - Value: Mix from black (0) to white (255)

02. Filtering

- Point processing** - $x_{ij} = f(p_{ij})$
 - Brightness** - $x_{ij} = p_{ij} + b$ (Clipping behavior)
 - Intensity scaling** - $x_{ij} = ap_{ij}$ (Increased/decreased contrast)
 - Normalization (Whitening)** - $x_{ij} = \frac{p_{ij} - \mu}{\sigma} \sigma^2 = \frac{\sum_{i=0}^I \sum_{j=0}^J (p_{ij} - \mu)^2}{IJ}$
 - Gamma** - $x_{ij} = 255(\frac{p_{ij}}{255})^\gamma$ where $\gamma > 0$
 - Intensity Histogram** - No data on location
 - Stretching** - $x = (p - f_{\min})(\frac{255}{f_{\max} - f_{\min}})$
 - Equalization** - Turn cumulative distribution linear
 - Get histogram: $h_k = \sum_{i=0}^I \sum_{j=0}^J \delta$ where $\delta = 1$ if $p_{ij} - k = 0$ and 0 otherwise
 - Estimate CDF: $c_k = \frac{\sum_{l=1}^k h_l}{IJ}$
 - Map $p_{ij} = k$ to new bin x_{ij} : $x_{ij} = 255 \text{CDF}(k)$
 - Pros: More contrast than normalization, less prone to outliers than stretching
 - Cons: More expensive
 - Good for foreground and background separation: Bi-modal histogram can be thresholded
 - Otsu's Method** - Automated thresholding
 - $T^* = \min_{T \in [0, 255]} (w_1(T)\sigma_1^2(T) + w_2(T)\sigma_2^2(T))$ where T^* is the optimal threshold that min. sum of weighted variances of object and background
 - $\sigma_1^2(T)$ and $\sigma_2^2(T)$: Variance of pixels less than or equal to and greater than threshold respectively
 - $w_1(T)$ and $w_2(T)$: Number of pixels less than or equal to and greater than threshold

• Correlation Filtering - Window (Moving average)

- Motivation: Reduce noise
 - Impulse noise: Random white occurrences of pixels
 - Salt and pepper noise: Rand. white and black pixels
 - Gaussian noise: Variations in intensity from Gaussian distribution $p_{ij} = \hat{p}_{ij} + \eta$ where $\eta \sim N(\mu_n, \sigma_n)$
- $x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$ where f is a **kernel** of weights
- Notation: $X = P \otimes F$
- Box Filter** - Blurs image

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Given kernel of width $2k + 1$, add padding: $\lfloor \frac{2k+1}{2} \rfloor$
 - Filling methods: Zero-padding, wrap around, copy edge, reflect
- Gaussian Filter** - Nearest neighboring pixels have more weight

$$f_{uv} = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

- Size of kernel: At some threshold, kernel will have lots of 0s, which is not useful
- Variance: Determines smoothing effect
- Template Matching** - Set template as kernel and match occurs at local max.
 - Normalized Cross-Correlation** - Normalizes using filter and input window
 - Motivation: Non-normalized is dominated by original image pixels
 - $x_{ij} = \frac{1}{|F||w_{ij}|} \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$ where $|\cdot|$ is magnitude of kernel (i.e. Square root of sum of squares)
- Convolution Filtering** - Flip kernel in both directions, then do cross-correlation
 - $x_{ij} = \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i-u, j-v}$
 - Notation: $X = F * P$
 - Convolution has nicer properties than correlation
 - Properties: Commutative, Associative, Distributive over addition, Scalar factor, Identity
- Non-Linear Filters** - Filters that perform non-linear operations (e.g. Median, Min., Max.)
 - Correlation and convolution are both linear operations

• Median Filtering


- Removes spikes: Good for removing impulse and salt and pepper noise
- No new pixel values introduced
- Edge preserving

03. Gradients

04. Lines

- Goal: Which points belong to which line?
- $y = mx + b$ $\frac{x}{a} + \frac{y}{b} = 1$
- $x \cos \theta + y \sin \theta = \rho$
- Line fitting** - $E = \frac{1}{N} \sum_i (y_i - (mx_i + b))^2$
 - Goal: Find parameters m and b that min. E
 - Using gradient descent: $b = \bar{y} - m\bar{x}$ $m = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$ where $\bar{y} = \frac{\sum_i y_i}{N}$ and $\bar{x} = \frac{\sum_i x_i}{N}$
 - Cons: Choice of error function, outliers

• Hough Transform - Parameterize shape and vote



- Initialize accumulator array $A(\theta, \rho) = 0$
- For each image edge point (x_i, y_i) :
 - For each θ :
 - Solve for $\rho = x_i \cos \theta + y_i \sin \theta$
 - $A(\theta, \rho) = A(\theta, \rho) + 1$
- Find maximum in $A(\theta, \rho)$
- Detected lines are given by $\rho^* = x \cos \theta^* + y \sin \theta^*$
- Circle** - $(x - a)^2 + (y - b)^2 = r^2$
 - If r is unknown, need to solve for 3 parameters (i.e. $A(a, b) \rightarrow A(a, b, r)$) and a point becomes a cone in 3D parameter space
 - Optimization: Use gradient
 - Use edge orientation ϕ_i to vote for 2 points, rather than whole circle
- Generalized Hough Transform** - Arbitrary shape
 - Given shape with boundary points p_i and reference point a
 - For each p_i , get displacement vector from a
 - Store in table with key ϕ_i and value (Displacement vectors)
 - For each edge pt., use ϕ_i to get vectors to vote for a
 - Application: Index with local pattern, instead of gradients
 - Tips: Soft voting, convert to edge image
 - Pros: Food with noise and occlusion
 - Cons: Complexity with num. of parameters, grid size

05. Segmentation

- Goal: Separate image into coherent regions
- Idea: **Clustering** - Group similar data points together
- Challenges: What makes 2 points same/different? Choice of features (e.g. Color, Intensity, Position), Which clustering algorithm?
- k-Means Clustering** - Iteratively re-assign points to nearest cluster center
 - Randomly initialize the cluster centers c_1, \dots, c_K
 - For each point p_i , find the closest c_j to put p_i in
 - Set c_j to be mean of points in cluster j
 - Repeat, if c_j have changed up to some threshold
 - Pros: Simple, Converges to local min.
 - Cons: Setting K , Sensitive to initial centers (Since k-means converges to local min.), Sensitive to outliers (Can add more clusters), Assumes spherical clusters
- Simple Linear Iterative Clustering (SLIC) Superpixels**
 - Superpixel** - Group of pixels that share common traits
 - Application: Inputs to other CV algo. since more compact representation with perceptual meaning
 - Num. of pixels: n_{tp} ; Target num. of superpixels: n_{sp}
 - Initial width of each superpixel: $s = \sqrt{n_{tp}/n_{sp}}$
 - Features: $z = [r, g, b, x, y]$
 - Color distance: $d_c = ||\langle r_j, g_j, b_j \rangle - \langle r_i, g_i, b_i \rangle||$
 - Spatial distance: $d_s = ||\langle x_j, y_j \rangle - \langle x_i, y_i \rangle||$
 - Scaling factors: d_{cm} and $d_{sm} = s$ set as max. expected values of d_c and d_s respectively
 - $D = \sqrt{(\frac{d_c}{d_{cm}})^2 + (\frac{d_s}{d_{sm}})^2} = \sqrt{d_c^2 + (\frac{d_s}{s})^2} c^2$
- Split img. into grid of size $s \times s$. Set cluster centers as lowest gradient position in 3×3 neighborhood from superpixel center to speed up convergence since initialize on value common to surrounding.
- For each cluster center, check distance to all pixels within $2s \times 2s$ neighborhood. Assign pixels to closest checked center.
- Update cluster centers using mean and repeat if not converged (Same as k-Means)
- Optional: Replace superpixel region by average value to create stained glass effect
- Modification of k-Means: Not random initialization, Compute pixel's distance only to closest set of cluster centers
- Mean-Shift Clustering** - Find local density maxima in feature space
 - Attraction basin** - Region in feature space for which all trajectories of centroids lead to same mode
 - Cluster** - All data points in attraction basin of a mode

1. For each data point:
 1. Define window around and get centroid
 2. Shift window to centroid
 3. Repeat until window centroid stops moving
- Segmentation with Mean Shift: Do mean shift and merge pixels in same attraction basin
- Choosing window size: Trial and error, Sample points and use avg. dist. to knn. (Num. of neighbors needs to be large enough to ensure increase in density)
 - Larger window size → Fewer clusters
- Pros: No assumptions on cluster shape, 1 parameter, Finds variable num. of modes (vs. specified k in k-Means), Robust to outliers
- Cons: Choosing h , Slow, Scales poorly with feature space dimension
- Optimizations:
 - After each run of mean shift, assign all points within radius r of end point to same cluster
 - Assign points in radius $c < r$ of search path to mode